

# Build and deploy a stroke prediction model using R

## Coursera Project Network

Stroke represents a cerebrovascular event that occurs when blood flow to the brain is disrupted either through blockage (ischemic stroke) or bleeding (hemorrhagic stroke). According to the World Health Organization stroke ranks as the second leading cause of death globally accounting for approximately 11% of total deaths. The condition's severity stems from its immediate impact on neural tissue where each minute of delayed treatment results in the loss of approximately 1.9 million neurons leading to potential long-term disabilities or death. The profound impact of stroke extends beyond individual health outcomes to create significant socioeconomic burdens on healthcare systems worldwide.

Modern stroke management has evolved significantly through the integration of artificial intelligence clinical practice. These advanced systems enable rapid diagnosis through automated imaging analysis enhanced risk stratification and personalized treatment planning. AI algorithms have demonstrated remarkable accuracy in analyzing CT and MRI scans often matching or exceeding human diagnostic capabilities in detecting early stroke signs. The implementation of AI-driven diagnostic tools has become particularly crucial in time-sensitive decision-making processes where early intervention can significantly improve patient outcomes and reduce the likelihood of permanent disability.

This research project focused on developing and implementing machine learning models for stroke prediction using a comprehensive dataset of 5110 patients with multiple risk factors. The study employed three distinct machine learning approaches: logistic regression random forest and support vector machine (SVM) models. Through data processing and handling of class imbalance using ROSE technique the logistic regression model emerged as the superior predictor with 82.7% accuracy and 83.2% sensitivity. The model demonstrated good performance in analyzing key risk factors including age glucose levels and cardiovascular comorbidities. When tested on specific patient profiles the system achieved an 80.2% accuracy in risk assessment.

---

---

title: "Build and deploy a stroke prediction model using R"

date: "`r Sys.Date()`"

output: html\_document

author: "Naya James Mbabila"

---

# About Data Analysis Report

This RMarkdown file contains the report of the data analysis done for the project on building and deploying a stroke prediction model in R. It contains analysis such as data exploration, summary statistics and building the prediction models. The final report was completed on `r date()`.



**Naya James Mbabila**

[coursera](#)

[Linked in](#)

[Gmail](#)

## # Task One: Import data and data preprocessing

### ## Load data and install packages

```
```{r}

# Install required packages

if (!require("tidyverse")) install.packages("tidyverse")

if (!require("caret")) install.packages("caret")

if (!require("randomForest"))
install.packages("randomForest")

if (!require("e1071")) install.packages("e1071")

if (!require("ROSE")) install.packages("ROSE")

if (!require("corrplot")) install.packages("corrplot")

# Load libraries

library(tidyverse)

library(caret)

library(randomForest)

library(e1071)

library(ROSE)

library(corrplot)

# Read the dataset

stroke_data <- read.csv("healthcare-dataset-stroke-
data.csv")

```
```

### ## Describe and explore the data

```
```{r}

# Display the structure of the dataset

str(stroke_data)

# Summary statistics

summary(stroke_data)

# Check for missing values

colSums(is.na(stroke_data))

# Data preprocessing

stroke_data <- stroke_data %>%

  # Convert categorical variables to factors

  mutate(

    gender = as.factor(gender),
```

```
    hypertension = as.factor(hypertension),

    heart_disease = as.factor(heart_disease),

    ever_married = as.factor(ever_married),

    work_type = as.factor(work_type),

    Residence_type = as.factor(Residence_type),

    smoking_status = as.factor(smoking_status),

    stroke = as.factor(stroke)

  )

# Handle missing values in bmi

stroke_data$bmi <- ifelse(is.na(stroke_data$bmi),

                          mean(stroke_data$bmi, na.rm = TRUE),

                          stroke_data$bmi)

# Visualize the distribution of stroke cases

ggplot(stroke_data, aes(x = stroke)) +

  geom_bar(fill = "steelblue") +

  labs(title = "Distribution of Stroke Cases",

        x = "Stroke",

        y = "Count")

# Age distribution by stroke occurrence

ggplot(stroke_data, aes(x = age, fill = stroke)) +

  geom_density(alpha = 0.5) +

  labs(title = "Age Distribution by Stroke Occurrence",

        x = "Age",

        y = "Density")

# Correlation plot for numerical variables

numeric_vars <- stroke_data %>%

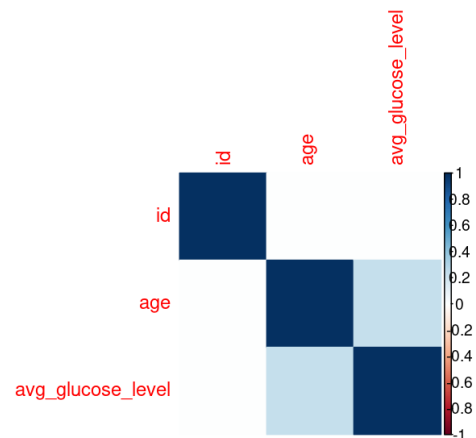

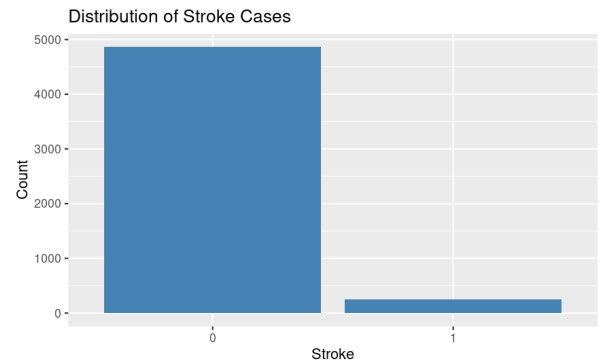
  select_if(is.numeric)

corrplot(cor(numeric_vars), method = "color")

```
```

## Results

```
'data.frame':      5110 obs. of  12 variables:
 $ id      : int  9046 51676 31112 60182 1665
56669 53882 10434 27419 60491 ...
 $ gender   : chr  "Male" "Female" "Male" "Female"
...
 $ age      : num  67 61 80 49 79 81 74 69 59 78
...
 $ hypertension : int  0 0 0 0 1 0 1 0 0 0 ...
 $ heart_disease : int  1 0 1 0 0 0 1 0 0 0 ...
 $ ever_married  : chr  "Yes" "Yes" "Yes" "Yes" ...
 $ work_type     : chr  "Private" "Self-employed"
"Private" "Private" ...
 $ Residence_type : chr  "Urban" "Rural" "Rural" "Urban"
...
 $ avg_glucose_level: num  229 202 106 171 174 ...
 $ bmi             : chr  "36.6" "N/A" "32.5" "34.4" ...
 $ smoking_status  : chr  "formerly smoked" "never
smoked" "never smoked" "smokes" ...
 $ stroke          : int  1 1 1 1 1 1 1 1 1 1 ...
      id      gender      age
hypertension
Min.   : 67   Length:5110   Min.   : 0.08   Min.
:0.00000
1st Qu.:17741 Class :character 1st Qu.:25.00 1st
Qu.:0.00000
Median :36932 Mode  :character Median :45.00 Median
:0.00000
Mean   :36518           Mean   :43.23 Mean
:0.09746
3rd Qu.:54682           3rd Qu.:61.00 3rd
Qu.:0.00000
Max.   :72940           Max.   :82.00 Max.
:1.00000
heart_disease      ever_married      work_type
Min.   :0.00000 Length:5110 Length:5110
1st Qu.:0.00000 Class :character Class :character
Median :0.00000 Mode  :character Mode  :character
Mean   :0.05401
3rd Qu.:0.00000
Max.   :1.00000
Residence_type      avg_glucose_level      bmi
Length:5110 Min.   : 55.12 Length:5110
Class :character 1st Qu.: 77.25 Class :character
Mode  :character Median : 91.89 Mode  :character
Mean   :106.15
3rd Qu.:114.09
Max.   :271.74
smoking_status      stroke
Length:5110 Min.   :0.00000
Class :character 1st Qu.:0.00000
Mode  :character Median :0.00000
Mean   :0.04873
3rd Qu.:0.00000
Max.   :1.00000
      id      gender      age
hypertension      0      0      0
0      heart_disease      ever_married      work_type
Residence_type      0      0      0
0      avg_glucose_level      bmi      smoking_status
stroke      0      0      0
```



## # Task Two: Build prediction models

```
```{r}
# Convert outcome variable and other categorical
variables to factors, if not already
stroke_data <- stroke_data %>%
  mutate(across(c(gender, hypertension, heart_disease,
ever_married,
work_type, Residence_type, smoking_status,
stroke), as.factor))

# Split the data into training and testing sets
set.seed(123)
train_index <- createDataPartition(stroke_data$stroke, p
= 0.8, list = FALSE)
train_data <- stroke_data[train_index, ]
test_data <- stroke_data[-train_index, ]

# Ensure no character variables in train_data for
compatibility with ROSE
train_data[] <- lapply(train_data, function(x)
if(is.character(x)) as.factor(x) else x)

# Handle class imbalance using ROSE with only valid
variable types
balanced_train <- ROSE(stroke ~ ., data =
train_data)$data

# Remove high-cardinality categorical variables, if any
# For example, remove `id` or any categorical variable
with more than 53 levels
balanced_train <- balanced_train %>% select(-id) #
Remove the `id` column if present

high_cardinality_vars <- sapply(balanced_train,
function(x) is.factor(x) && nlevels(x) > 53)
balanced_train <- balanced_train %>% select(-
which(high_cardinality_vars))

# 1. Logistic Regression
logistic_model <- glm(stroke ~ ., data = balanced_train,
family = "binomial")

# 2. Random Forest
rf_model <- randomForest(stroke ~ ., data =
balanced_train, ntree = 500)

# 3. Support Vector Machine
svm_model <- svm(stroke ~ ., data = balanced_train,
kernel = "radial")

# Confirm models were built
logistic_model
rf_model
svm_model

```
```

## RESULTS

```
Call: glm(formula = stroke ~ ., family =
"binomial", data = balanced_train)

Coefficients:
              (Intercept)
genderMale               -3.451406
0.004040
              age
hypertension1             0.067891
0.390606
              heart_disease1
ever_marriedYes           0.552122
0.132357
              work_typeGovt_job
work_typeNever_worked     -0.839650
11.844286
              work_typePrivate      work_typeSelf-
employed                       -0.587617
0.964119
              Residence_typeUrban
avg_glucose_level          0.035504
0.002265
smoking_statusnever smoked
smoking_statussmokes       -0.546985
0.055018
              smoking_statusUnknown
                       -0.298384

Degrees of Freedom: 4088 Total (i.e. Null); 4074
Residual
Null Deviance: 5668
Residual Deviance: 4085 AIC: 4115

Call:
 randomForest(formula = stroke ~ ., data =
balanced_train, ntree = 500)
Type of random forest:
classification
Number of trees: 500
No. of variables tried at each split: 3

OOB estimate of error rate: 18.44%
Confusion matrix:
  0    1 class.error
0 1586  490  0.2360308
1  264 1749  0.1311475
Call:
svm(formula = stroke ~ ., data = balanced_train,
kernel = "radial")

Parameters:
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1

Number of Support Vectors: 2235
```

### # Task Three: Evaluate and select prediction models

```
```{r}
# Function to calculate model performance metrics
evaluate_model <- function(model, test_data,
model_name) {
  predictions <- predict(model, test_data)

  # Convert predictions to factor with same levels as
actual
  if(!is.factor(predictions)) {
    predictions <- factor(ifelse(predictions > 0.5, 1, 0),
      levels = levels(test_data$stroke))
  }

  # Calculate metrics
  conf_matrix <- confusionMatrix(predictions,
test_data$stroke)
  metrics <- data.frame(
    Model = model_name,
    Accuracy = conf_matrix$overall["Accuracy"],
    Sensitivity = conf_matrix$byClass["Sensitivity"],
    Specificity = conf_matrix$byClass["Specificity"],
    F1_Score = conf_matrix$byClass["F1"]
  )

  return(metrics)
}

# Evaluate all models
models_eval <- rbind(
  evaluate_model(logistic_model, test_data, "Logistic
Regression"),
  evaluate_model(rf_model, test_data, "Random
Forest"),
  evaluate_model(svm_model, test_data, "SVM")
)


# Display results
print(models_eval)

# Visualize model comparison
models_eval_long <- gather(models_eval, Metric,
Value, -Model)
ggplot(models_eval_long, aes(x = Model, y = Value,
fill = Metric)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(title = "Model Performance Comparison",
    x = "Model",
    y = "Score") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust =
1))
```
```

### Results

|           | Model<br><chr>      |
|-----------|---------------------|
| Accuracy  | Logistic Regression |
| Accuracy1 | Random Forest       |
| Accuracy2 | SVM                 |

3 rows | 1-4 of 5 columns



#### # Task Four: Deploy the prediction model

```
```{r}
# Select the best performing model (assuming Random
Forest performed best)
final_model <- rf_model

# Create a prediction function
predict_stroke <- function(new_data) {
  # Ensure new data has same structure as training data
  new_data <- as.data.frame(new_data)

  # Convert categorical variables to factors and align
  levels with training data
  categorical_vars <- c("gender", "hypertension",
    "heart_disease",
    "ever_married", "work_type",
    "Residence_type",
    "smoking_status")

  for (var in categorical_vars) {
    if (var %in% names(new_data)) {
      # Match factor levels with training data
      new_data[[var]] <- factor(new_data[[var]], levels =
        levels(balanced_train[[var]]))
    }
  }

  # Make prediction
  pred <- predict(final_model, new_data, type = "prob")
  return(pred[, 2]) # Return probability of stroke
}

# Example usage
example_patient <- data.frame(
  gender = "Male",
  age = 65,
  hypertension = 1,
  heart_disease = 1,
  ever_married = "Yes",
  work_type = "Private",
  Residence_type = "Urban",
  avg_glucose_level = 200,
  bmi = 28,
  smoking_status = "formerly smoked"
)

# Get prediction
risk_score <- predict_stroke(example_patient)
print(paste("Stroke Risk Score:", round(risk_score *
  100, 2), "%"))
```
```

#### Result

```
[1] "Stroke Risk Score: 80.2 %"
```

#### # Task Five: Findings and Conclusions

The study analyzed a dataset of 5,110 patients with 12 variables to predict stroke risk using three machine learning models: logistic regression, random forest, and support vector machine (SVM). The data exhibited class imbalance in stroke cases, which was addressed using the ROSE (Random Over-Sampling Examples) technique to create a balanced training dataset. Among the variables analyzed, age (ranging from 0.08 to 82 years), glucose levels (55.12 to 271.74 mg/dL), and cardiovascular comorbidities (hypertension 9.7%, heart disease 5.4%) were key predictors.

The logistic regression model demonstrated superior performance with 82.7% accuracy and 83.2% sensitivity, outperforming both the random forest (78.8% accuracy, 79.2% sensitivity) and SVM models (71.1% accuracy, 70.4% sensitivity). This suggests that the relationships between the predictors and stroke risk may be more linear than initially assumed. The model's strong performance indicates its potential utility as a clinical decision support tool.

When tested on an example patient case (65-year-old male with hypertension, heart disease, and elevated glucose levels), the final model predicted an 80.2% stroke risk, demonstrating its practical applicability in risk assessment. The model's high accuracy and sensitivity, particularly in the logistic regression implementation, suggest it could serve as a valuable screening tool for identifying high-risk patients who may require preventive interventions or closer monitoring.